

Daily Macro Brief Agent - 1-Page Memo

Design Tradeoffs

This prototype is designed to run manually or through a scheduler on an always-on machine. I chose *Gmail SMTP* rather than *Telegram* because email fits the intended daily workflow and requires no new user habit. *Gemini 2.5 Flash-Lite* is the default LLM because the model's job is narrow at this stage: synthesize structured inputs into a concise macro brief. Code, rather than the LLM, owns market data, tables, charting, source links, validation, logging, and delivery.

The core design choice is to separate facts from writing. Market numbers come from public APIs or scraping: *Yahoo Finance chart data*, *Japan MOF JGB data*, *Frankfurter FX reference rates*, and *CoinGecko BTC data*. If live market data is unavailable and no cached data exists, the value cells are left blank intentionally. Calendar rows come from the *Fair Economy / Forex Factory weekly feed*; RSS sources feed *Theme Radar*. The LLM receives only structured facts and is validated for JSON shape, word limits, market-number consistency, source reuse, and portfolio logic before output is accepted.

The brief is tailored to a macro PM rather than a news reader. The dashboard is a table with "Reading". The 3 Things section is exactly three items with portfolio implications. The chart is *USD/JPY* because the assumed book is long *USD/JPY*; the current chart uses about three months of history and highlights the latest five observations. *Theme Radar* uses curated RSS sources with source-depth labels, so readers can see when the agent is using feed-level text rather than full article text. Scheduling was tested: *GitHub manual send* worked, *GitHub scheduled triggers* were unreliable in short-window tests, and a *MacBook launchd* scheduled send succeeded end to end.

Position And Theme Assumptions

Because the PDF does not reveal the real book, the prototype uses explicit assumptions stored in `inputs/portfolio/positions.csv`. The assumed book is long *USD/JPY*, overweight *gold*, and exposed to an *EM debt basket*. The assumed house themes are *higher-for-longer rates*, *fiscal pressure and term premium*, *China demand weakness*, and *policy divergence* across major central banks. Position rows are effective-date updates, so the latest prior row carries forward until changed or closed.

Three V2 Features Not Completed

1. Add redundant or paid *market/calendar providers* so the agent is less dependent on free public feeds.
2. Expand *Theme Radar* from RSS-level text to permission-aware full-text reading for research notes, speeches, and selected newsletters. At that point, a more advanced LLM may be more useful.
3. Use feedback history and richer book information to improve ranking, prompt construction, and brief personalization over time.

One-Month Full-Time Roadmap

With one month full-time, I would turn this into a production briefing system: add robust provider redundancy, source-health monitoring, schedule alerting, richer article ingestion, feedback-driven ranking, and a cleaner deployment target such as a VPS or cloud scheduler. I would also add a daily evaluation report showing source freshness, failed fetches, token use, runtime, and whether the delivered brief met all formatting and factual guardrails.

Actual Hours Spent

Building the agent was relatively fast; I spent about 2-3 hours on the initial scaffold. I then spent another full day, more than 10 additional hours, checking that the content was real, updated, reader-friendly, well-structured, and easier to extend later.